

WORKSHEET - 4 (AI)

Problem 1:

1.

```
[2]: import pandas as pd
      import numpy as np
      import matplotlib.pyplot as plt
      import time
```

Problem - 1: Perform a classification task with knn from scratch.



1. Load the Dataset:
 - Read the dataset into a pandas DataFrame.
 - Display the first few rows and perform exploratory data analysis (EDA) to understand the dataset (e.g., check data types, missing values, summary statistics).

```
[3]: df = pd.read_csv("/content/drive/MyDrive/Concept and Technology Of AI/Week4/diabetes_.csv")
print("\nFirst 5 rows:")
print(df.head())

print("\nData types: ")
print(df.dtypes)

print("\nMissing values: ")
print(df.isna().sum())

print("\nSummary statistics: ")
print(df.describe(include='all'))
```

First 5 rows:

```
Pregnancies Glucose BloodPressure SkinThickness Insulin BMI \
0           6      148            72          35       0   33.6
1           1      85             66          29       0   26.6
2           8      183            64           0       0   23.3
3           1      89             66          23     94   28.1
4           0      137            40          35    168  43.1
```

```
DiabetesPedigreeFunction Age Outcome
0           0.627      50      1
1           0.351      31      0
2           0.672      32      1
3           0.167      21      0
4           2.288      33      1
```

Data types:

```
Pregnancies          int64
Glucose              int64
BloodPressure        int64
SkinThickness        int64
Insulin              int64
BMI                 float64
DiabetesPedigreeFunction float64
Age                 int64
Outcome              int64
dtype: object
```

Summary statistics:

```
Pregnancies Glucose BloodPressure SkinThickness Insulin \
count    768.000000 768.000000 768.000000 768.000000 768.000000
mean     3.845052 120.894531 69.105469 20.536458 79.799479
std      3.369578 31.972618 19.355807 15.952218 115.244002
min      0.000000 0.000000 0.000000 0.000000 0.000000
25%     1.000000 99.000000 62.000000 0.000000 0.000000
50%     3.000000 117.000000 72.000000 23.000000 30.500000
75%     6.000000 140.250000 80.000000 32.000000 127.250000
max     17.000000 199.000000 122.000000 99.000000 846.000000
```

```
BMI DiabetesPedigreeFunction Age Outcome
count 768.000000 768.000000 768.000000 768.000000
mean  31.992578 0.471876 33.240885 0.348958
std   7.884160 0.331329 11.760232 0.476951
min   0.000000 0.078000 21.000000 0.000000
25%  27.300000 0.243750 24.000000 0.000000
50%  32.000000 0.372500 29.000000 0.000000
75%  36.600000 0.626250 41.000000 1.000000
max  67.100000 2.420000 81.000000 1.000000
```

2.

- Handle Missing Data:
 - Handle any missing values appropriately, either by dropping or imputing them based on the data.

```
[13]: df_handle = df.copy()
for col in df_handle.columns:
    if df_handle[col].isna().any():
        median_value = df_handle[col].median() if np.issubdtype(df_handle[col].dtype, np.number) else df_handle[col].mode()[0]
        df_handle[col] = df_handle[col].fillna(median_value)

print("Any remaining NaNs?", df_handle.isna().any().any())
Any remaining NaNs? False
```

3.

- Feature Engineering:
 - Separate the feature matrix (X) and target variable (y).
 - Perform a train - test split from scratch using a 70% – 30% ratio.

```
[11]: X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values

np.random.seed(42)
indices = np.random.permutation(len(X))

train_size = int(0.7 * len(X))
train_idx = indices[:train_size]
test_idx = indices[train_size:]

X_train = X[train_idx]
X_test = X[test_idx]
y_train = y[train_idx]
y_test = y[test_idx]
print("Train size:", X_train.shape[0], "Test size:", X_test.shape[0])
```

Train size: 537 Test size: 231

4.

4. Implement KNN:

- Build the KNN algorithm from scratch (no libraries like scikit-learn for KNN).
- Compute distances using Euclidean distance.
- Write functions for:
 - Predicting the class for a single query.
 - Predicting classes for all test samples.
- Evaluate the performance using accuracy.

```
: def euclidean_distance(a: np.ndarray, b: np.ndarray):
    return np.sqrt(np.sum(a - b) ** 2)

def knn_predict_one(x_query: np.ndarray, X_train: np.ndarray, y_train: np.ndarray, k: int = 5):
    dists = np.sqrt(np.sum((X_train - x_query) ** 2, axis=1))
    nn_idx = np.argpartition(dists, k)[:k]
    nn_labels = y_train[nn_idx]
    values, counts = np.unique(nn_labels, return_counts=True)
    return values[np.argmax(counts)]

def knn_predict(X_query: np.ndarray, X_train: np.ndarray, y_train: np.ndarray, k: int = 5):
    preds = np.zeros(X_query.shape[0], dtype=y_train.dtype)
    for i in range(X_query.shape[0]):
        preds[i] = knn_predict_one(X_query[i], X_train, y_train, k)
    return preds

def accuracy_score(y_true: np.ndarray, y_pred: np.ndarray):
    return (y_true == y_pred).mean()

k_baseline = 5
t0 = time.time()
y_pred_baseline = knn_predict(X_test, X_train, y_train, k=k_baseline)
t1 = time.time()

acc_baseline = accuracy_score(y_test, y_pred_baseline)
print(f"Baseline (unscaled) KNN | k={k_baseline}: Accuracy={acc_baseline:.4f}, Time={t1 - t0:.4f}s")
Baseline (unscaled) KNN | k=5: Accuracy=0.7143, Time=0.0255s
```

Problem - 2 - Experimentation

1.

Problem - 2 - Experimentation:

1. Repeat the Classification Task:

- Scale the Feature matrix X.
- Use the scaled data for training and testing the kNN Classifier.
- Record the results.

```
[20]: X_mean = X_train.mean(axis=0)
X_std = X_train.std(axis=0)
X_std[X_std == 0] = 1.0

X_train_scaled = (X_train - X_mean) / X_std
X_test_scaled = (X_test - X_mean) / X_std

k_scaled = 3
t0 = time.time()
y_pred_scaled = knn_predict(X_test_scaled, X_train_scaled, y_train, k=k_scaled)
t1 = time.time()

acc_scaled = accuracy_score(y_test, y_pred_scaled)
print(f"Scaled KNN | k={k_scaled}: Accuracy={acc_scaled:.4f}, Time={(t1 - t0):.4f}s")

Scaled KNN | k=3: Accuracy=0.6970, Time=0.0222s
```

2.

2. Comparative Analysis: Compare the Results -

- Compare the accuracy and performance of the kNN model on the original dataset from problem 1 versus the scaled dataset.
- Discuss:
 - How scaling impacted the KNN performance.
 - The reason for any observed changes in accuracy.

```
[22]: print("Comparison at k=3")
print(f"- Unscaled accuracy: {acc_baseline:.4f}")
print(f"- Scaled accuracy: {acc_scaled:.4f}")
print("Observation: Scaling changes feature magnitudes, which can impact neighbor selection and thus accuracy.")

Comparison at k=3
- Unscaled accuracy: 0.7143
- Scaled accuracy: 0.6970
Observation: Scaling changes feature magnitudes, which can impact neighbor selection and thus accuracy.
```

3.

Problem - 3 - Experimentation with k:

1. Vary the number of neighbors - k:

- Run the KNN model on both the original and scaled datasets for a range of:

k= 1, 2, 3, . . . 15

- For each k, record:

- Accuracy.
 - Time taken to make predictions.

```
[24]: ks = list(range(1, 16))
acc_unscaled_list, time_unscaled_list = [], []
acc_scaled_list, time_scaled_list = [], []

for k in ks:
    t0 = time.time()
    y_pred_u = knn_predict(X_test, X_train, y_train, k=k)
    t1 = time.time()
    acc_unscaled_list.append(accuracy_score(y_test, y_pred_u))
    time_unscaled_list.append(t1 - t0)

    t0 = time.time()
    y_pred_s = knn_predict(X_test_scaled, X_train_scaled, y_train, k=k)
    t1 = time.time()
    acc_scaled_list.append(accuracy_score(y_test, y_pred_s))
    time_scaled_list.append(t1 - t0)

best_k_unscaled = ks[int(np.argmax(acc_unscaled_list))]
best_k_scaled = ks[int(np.argmax(acc_scaled_list))]
```

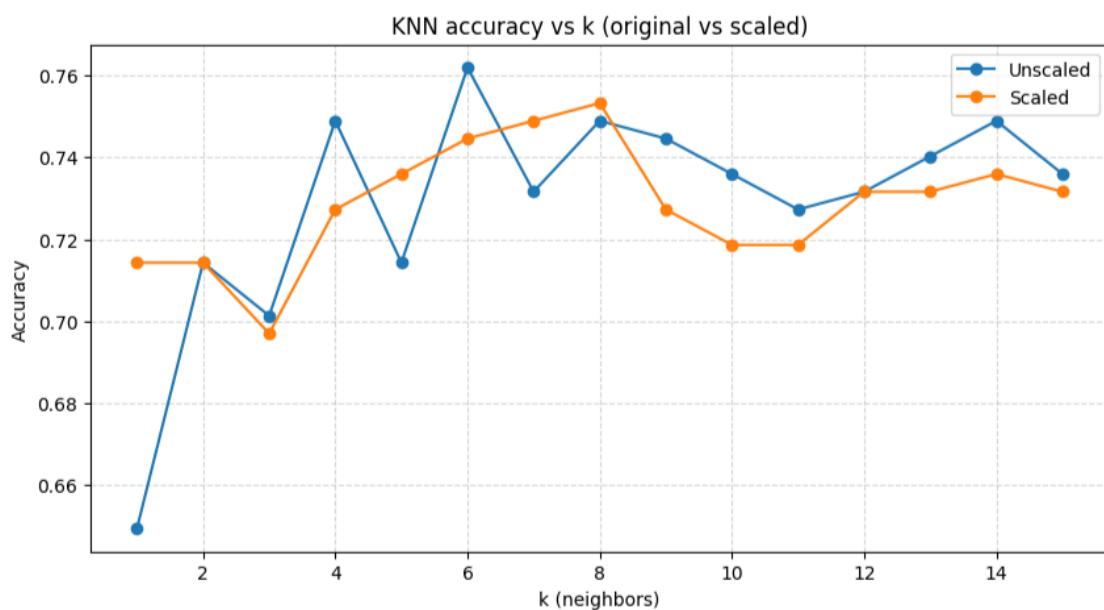
2. Visualize the Results

2. Visualize the Results:

- Plot the following graphs:
 - k vs. Accuracy for original and scaled datasets.
 - k vs. Time Taken for original and scaled datasets.

```
[25]: plt.figure(figsize=(10, 5))
plt.plot(ks, acc_unscaled_list, marker='o', label='Unscaled')
plt.plot(ks, acc_scaled_list, marker='o', label='Scaled')
plt.xlabel('k (neighbors)')
plt.ylabel('Accuracy')
plt.title('KNN accuracy vs k (original vs scaled)')
plt.grid(True, linestyle='--', alpha=0.5)
plt.legend()
plt.show()

print(f"Best k (unscaled): {best_k_unscaled} with accuracy {max(acc_unscaled_list):.4f}")
print(f"Best k (scaled): {best_k_scaled} with accuracy {max(acc_scaled_list):.4f}")
```



```
Best k (unscaled): 6 with accuracy 0.7619
Best k (scaled): 8 with accuracy 0.7532
```

3. Analyze and Discuss

3. Analyze and Discuss:

- Discuss how the choice of k affects the accuracy and computational cost.
- Identify the optimal k based on your analysis.

```
[26]: plt.figure(figsize=(10, 5))
plt.plot(ks, time_unscaled_list, marker='o', label='Unscaled')
plt.plot(ks, time_scaled_list, marker='o', label='Scaled')
plt.xlabel('k (neighbors)')
plt.ylabel('Prediction time (s)')
plt.title('KNN prediction time vs k (original vs scaled)')
plt.grid(True, linestyle='--', alpha=0.5)
plt.legend()
plt.show()
```

