

To - Do - 1 :

#2

First 5 rows:

	Math	Reading	Writing
0	48	68	63
1	62	81	72
2	79	80	78
3	76	83	79
4	59	64	62

Last 5 rows:

	Math	Reading	Writing
995	72	74	70
996	73	86	90
997	89	87	94
998	83	82	78
999	66	66	72

#3

Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 3 columns):
 #   Column   Non-Null Count  Dtype  
--- 
 0   Math      1000 non-null   int64  
 1   Reading   1000 non-null   int64  
 2   Writing   1000 non-null   int64  
dtypes: int64(3)
memory usage: 23.6 KB
```

#4

Descriptive info:

	Math	Reading	Writing
count	1000.000000	1000.000000	1000.000000
mean	67.290000	69.872000	68.616000
std	15.085008	14.657027	15.241287
min	13.000000	19.000000	14.000000
25%	58.000000	60.750000	58.000000
50%	68.000000	70.000000	69.500000
75%	78.000000	81.000000	79.000000
max	100.000000	100.000000	100.000000

To – Do – 2 :

```
X shape: (2, 1000)
W shape: (2, 1)
Y shape: (1000,)
Y_pred shape: (1000, 1)
```

To – Do – 5 :

```
Proceed Further
Cost function output: 0.0
```

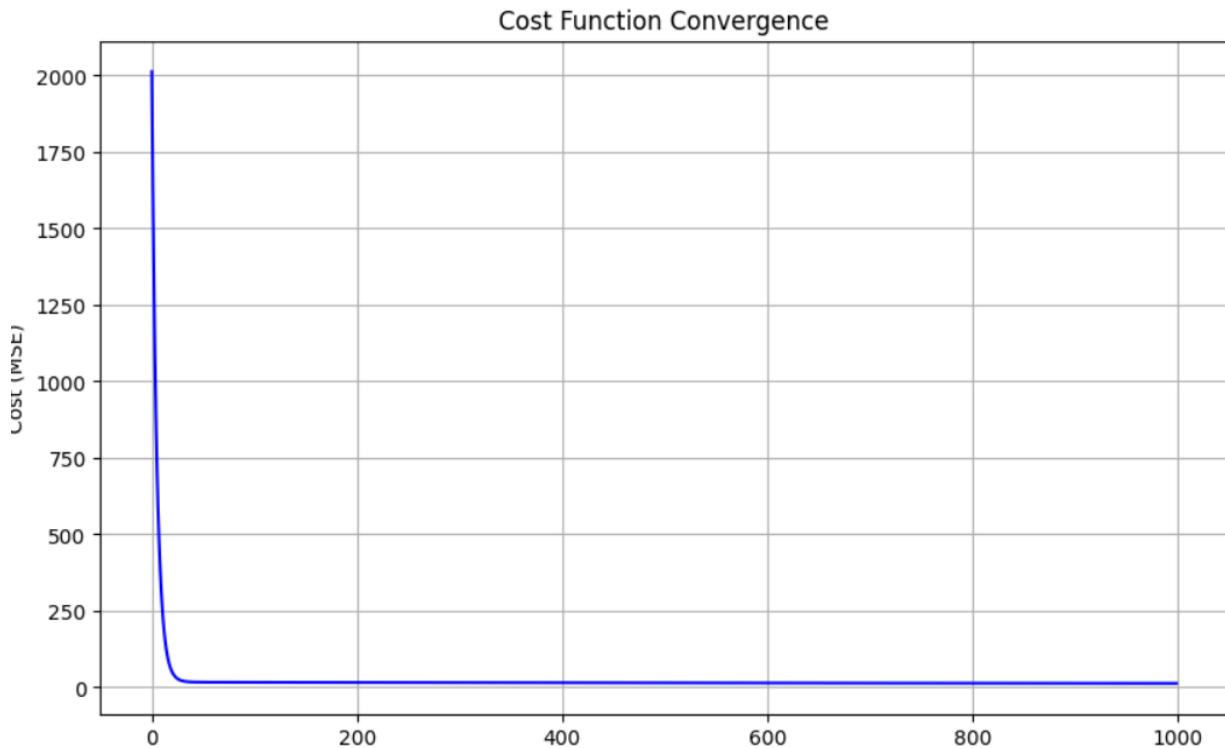
To – Do – 7 :

To – Do – 10 :

```
Final Weights: [[0.34811659]
[0.64614558]]
Cost History (First 10 iterations): [np.float64(2013.165570783755), np.float64(1640.286832599692), np.float64(1337.0619994901588), np.float64(1090.4794
892850578), np.float64(889.9583270083234), np.float64(726.8940993009545), np.float64(594.2897260808594), np.float64(486.4552052951635), np.float64(398.
7634463599484), np.float64(327.4517147324688)]
RMSE on Test Set: 5.2798239764188635
R-Squared on Test Set: 0.8886354462786421
```

To – Do – 11 :

1.



The model's performance is acceptable and the cost function decreases smoothly indicating proper convergence. The RMSE value being around 5.28 shows that the prediction error is low and R-squared value of approx 0.89 indicates that the model explains most of the variance in writing marks. Therefore, the model neither underfits nor overfits the data.

2.

```
Learning rate: 1e-06
Final Cost: 16.535602355147176
RMSE: 5.856694748793876
R2: 0.8629707528684534
```

```
Learning rate: 1e-05
Final Cost: 13.150619992105618
RMSE: 5.2798239764188635
R2: 0.8886354462786421
```

```
Learning rate: 0.0001
Final Cost: 10.26076310841341
RMSE: 4.792607360540954
R2: 0.908240340333986
```

According to the data above, as the learning rate increases from 0.000001 to 0.0001, the model shows improved convergence with decrease in final cost and RMSE and increase in R-Squared. Therefore, smaller learning rate results in slower convergence while moderate learning rate results in quicker convergence and is more effective but with bigger learning rate it might start to diverge.